

multiprocessing

importing required libraries and our shared library

```
In [1]: import ctypes
import multiprocessing
import os
import time
```

```
In [2]: _libInC = ctypes.CDLL('./libMyLib.so')
```

Here, we slightly adjust our Python wrapper to calculate the results and print it. There is also some additional casting to ensure that the result of the `libInC.myAdd()` is an int32 type.

```
In [3]: def addC_print(_i, a, b, time_started):
    val = ctypes.c_int32(_libInC.myAdd(a, b)).value #cast the result to a 32 bit integer
    end_time = time.time()
    print('CPU_{} Add: {} in {}'.format(_i, val, end_time - time_started))

def multC_print(_i, a, b, time_started):
    val = ctypes.c_int32(_libInC.myMultiply(a, b)).value #cast the result to a 32 bit integer
    end_time = time.time()
    print('CPU_{} Multiply: {} in {}'.format(_i, val, end_time - time_started))
```

Now for the fun stuff.

The multiprocessing library allows us to run simultaneous code by utilizing multiple processes. These processes are handled in separate memory spaces and are not restricted to the Global Interpreter Lock (GIL).

Here we define two processes, one to run the `_addCprint` and another to run the `_multCprint()` wrappers.

Next we assign each process to be run on different CPUs

```
In [4]: procs = [] # a future list of all our processes

# Launch process1 on CPU0
p1_start = time.time()
p1 = multiprocessing.Process(target=addC_print, args=(0, 3, 5, p1_start)) # the first
os.system("taskset -p -c {} {}".format(0, p1.pid)) # taskset is an os command to pin the
p1.start() # start the process
procs.append(p1)

# Launch process2 on CPU1
p2_start = time.time()
p2 = multiprocessing.Process(target=multC_print, args=(1, 3, 5, p2_start)) # the first
os.system("taskset -p -c {} {}".format(1, p2.pid)) # taskset is an os command to pin the
p2.start() # start the process
procs.append(p2)
```

```

p1Name = p1.name # get process1 name
p2Name = p2.name # get process2 name

# Here we wait for process1 to finish then wait for process2 to finish
p1.join() # wait for process1 to finish
print('Process 1 with name, {}, is finished'.format(p1Name))

p2.join() # wait for process2 to finish
print('Process 2 with name, {}, is finished'.format(p2Name))

```

```

taskset: invalid PID argument: 'None'
taskset: invalid PID argument: 'None'
CPU_1 Multiply: 15 in 2.065248489379883
CPU_0 Add: 8 in 4.136515378952026
Process 1 with name, Process-1, is finished
Process 2 with name, Process-2, is finished

```

Return to 'main.c' and change the amount of sleep time (in seconds) of each function.

For different values of sleep(), explain the difference between the results of the 'Add' and 'Multiply' functions and when the Processes are finished.

Lab work

One way around the GIL in order to share memory objects is to use multiprocessing objects. Here, we're going to do the following.

1. Create a multiprocessing array object with 2 entries of integer type.
2. Launch 1 process to compute addition and 1 process to compute multiplication.
3. Assign the results to separate positions in the array.
 - A. Process 1 (add) is stored in index 0 of the array (array[0])
 - B. Process 2 (mult) is stored in index 1 of the array (array[1])
4. Print the results from the array.

Thus, the multiprocessing Array object exists in a *shared memory* space so both processes can access it.

Array documentation:

<https://docs.python.org/2/library/multiprocessing.html#multiprocessing.Array>

typecodes/types for Array:

'c': ctypes.c_char

'b': ctypes.c_byte

'B': ctypes.c_ubyte
 'h': ctypes.c_short
 'H': ctypes.c_ushort
 'i': ctypes.c_int
 'I': ctypes.c_uint
 'l': ctypes.c_long
 'L': ctypes.c_ulong
 'f': ctypes.c_float
 'd': ctypes.c_double

Try to find an example

You can use online resources to find an example for how to use multiprocessing Array

```
In [13]: def addC_no_print(_i, a, b, returnValues):
    ...
    Params:
        _i : Index of the process being run (0 or 1)
        a, b : Integers to add
        returnValues : Multiprocessing array in which we will store the result at index
    ...
    val = ctypes.c_int32(_libInC.myAdd(a, b)).value
    returnValues[0] = val
    # TODO: add code here to pass val to correct position returnValues

def multC_no_print(_i, a, b, returnValues):
    ...
    Params:
        _i : Index of the process being run (0 or 1)
        a, b : Integers to multiply
        returnValues : Multiprocessing array in which we will store the result at index
    ...
    val = ctypes.c_int32(_libInC.myMultiply(a, b)).value
    returnValues[1] = val
    # TODO: add code here to pass val to correct position of returnValues

procs = []

# TODO: define returnValues here. Check the multiprocessing docs to see
# about initializing an array object for 2 processes.
# Note the data type that will be stored in the array
returnValues = multiprocessing.Array('i', range(2))

p1 = multiprocessing.Process(target=addC_no_print, args=(0, 3, 5, returnValues)) # the
os.system("taskset -p -c {} {}".format(0, p1.pid)) # taskset is an os command to pin t
p1.start() # start the process
```

```
procs.append(p1)

p2 = multiprocessing.Process(target=multC_no_print, args=(1, 3, 5, returnValues)) # this
os.system("taskset -p -c {} {}".format(1, p2.pid)) # taskset is an os command to pin t
p2.start() # start the process
procs.append(p2)

# Wait for the processes to finish
for p in procs:
    pName = p.name # get process name
    p.join() # wait for the process to finish
    print('{} is finished'.format(pName))

# TODO print the results that have been stored in returnValues
print(returnValues[:])
```

```
taskset: invalid PID argument: 'None'
taskset: invalid PID argument: 'None'
invalid PID argument: 'None'
Process-17 is finished
Process-18 is finished
[8, 15]
```